



Dynamic Linker

Sooner or later you'll reach the point where you want shared libraries. Here we won't discuss the difference between static and dynamic linking, you should be already familiar with that.

This article is about ELF on x86_64 architecture, but can be easily adopted to other systems as the concepts are similar.

Contents

Home work

Memory Layout

Segment Local Calls

Inter-segment Calls

Implementing a dynamic linker

Locating the GOT

What's in the GOT?

Where are my libraries?

Symbol look up

Gimme code!

Summary

Home work

Memory Layout

I suppose you have created a new, empty address space, and you've already loaded the executable in it. I also assume that you've loaded the libc shared library after the executable's segment. And now you're stuck, because you don't know how to call printf from the executable.

As your executable was loaded by you, you should know the virtual address of ELF magic bytes in the memory. Use that for start.

```
Elf64_Ehdr *ehdr = (Elf64_Ehdr *) (ptr);

if (!memcmp(ehdr->e_ident, ELF_MAGIC, ELF_MAGIC) &&
    ehdr->e_ident[EI_CLASS] == ELFCLASS64 &&
    ehdr->e_ident[EI_DATA] == ELFDATA2LSB &&
    ehdr->e_type == ET_EXEC && ehdr->e_shnum > 0)
{
    // We have a valid image with sections
}
```

First we check the magic bytes and the format of the ELF64. As an extra, we also check whether it's executable and has a non-empty section table (we'll going to need it).

Now let's see what the GNU toolchain does for us to find printf.

Segment Local Calls

In order to figure that out, first we should know how a segment local call works. For that we'll use a very minimal source.

```
void localfunction()
{
}

int main(int, char**)
{
    localfunction();
}
```

That compiles to:

```
$ objdump -d test
00000000020016f <localfunction>:
 20016f: 55                push    %rbp
 200170: 48 89 e5          mov     %rsp,%rbp
 200173: 90                nop
 200174: 5d                pop     %rbp
 200175: c3                retq

000000000200176 <main>:
 200176: 55                push    %rbp
 200177: 48 89 e5          mov     %rsp,%rbp
 20017a: b8 00 00 00 00    mov     $0x0,%eax
 20017f: e8 eb ff ff ff    callq   20016f <localfunction>
 200184: 90                nop
 200185: 5d                pop     %rbp
 200186: c3                retq
```

That's trivial, a rip relative addressing is used at 20017f.

Inter-segment Calls

Now let's modify the source a bit to use a libc call:

```
int main(int, char**)
{
    printf("Hello World");
}
```

Compile and see what's generated.

```
00000000020016f <main>:
 20016f: 55                push    %rbp
 200170: 48 89 e5          mov     %rsp,%rbp
 200173: 48 8d 3d b6 01 00 00    lea     0x1b6(%rip),%rdi    # 200330 <_DYNAMIC+0x110>
 20017a: b8 00 00 00 00    mov     $0x0,%eax
 20017f: e8 8c 00 00 00    callq   200210 <printf@plt>
 200184: 90                nop
 200185: 5d                pop     %rbp
 200186: c3                retq

000000000200200 <printf@plt-0x10>:
 200200: ff 35 4a 0e 00 00    pushq   0xe4a(%rip)        # 201050 <_GLOBAL_OFFSET_TABLE_+0x8>
 200206: ff 25 4c 0e 00 00    jmpq    *0xe4c(%rip)      # 201058 <_GLOBAL_OFFSET_TABLE_+0x10>
 20020c: 0f 1f 40 00        nopl    0x0(%rax)
```

```

0000000000200210 <printf@plt>:
200210: ff 25 4a 0e 00 00    jmpq    *0xe4a(%rip)          # 201060 <_GLOBAL_OFFSET_TABLE_+0x18>
200216: 68 00 00 00 00      pushq   $0x0
20021b: e9 e0 ff ff ff      jmpq    200200 <main+0x91>

```

What? Two more local functions? What happened here? The GNU toolchain has a concept for lazy run-time linking. That means the address is not resolved until it's referenced. To achieve that, it needs helper functions (generated to the .plt section in the text segment).

When the CPU executes this code, it will first call a normal local function at 20017f. That function is one of the helpers and its purpose is to load an address from GOT+0x18 and jump to it. By default, the value points to the next instruction, which saves 0 to stack and calls the other helper function at 200200. That one is the reference resolver, and its job is to replace the address in GOT with a relocated address.

Because the resolver function is not known at link time, its address is also in GOT at 0x10. What's more it can receive one argument, stored at GOT+0x8. We can also spot that at 200206 the instruction is a jump and not a call, so resolver never returns to this helper, instead it should jump to the relocated address.

Implementing a dynamic linker

All that disassembly teach us two things about the dynamic linker:

1. it has to locate and write the GOT
2. it has two parts: load time linker and a run time resolver.

The first part runs before the thread is started and saves the second part's address and argument into GOT. On the other hand, the second part runs when the thread is already running, and saves relocated addresses into GOT. As you can see, both parts require the address of GOT. To save resources, one should not locate the GOT twice: this is where the resolver's argument came in. While you can technically insert anything into this slot in the GOT, your best bet will be to store something that will easily identify the requesting program — for instance, the address of the start of the GOT.

To proceed we'll have to locate the GOT in memory and figure out what entries it has.

Locating the GOT

Time to peek on what's in the object file.

```

$ readelf -a test
Section Headers:
 [Nr] Name                Type              Address            Offset
     Size              EntSize          Flags  Link  Info  Align
 [10] '.got.plt'            PROGBITS          000000000201048    00001048
     0000000000000020  0000000000000008  WA           0      0      8

Symbol table '.symtab' contains 23 entries:
Num:  Value              Size Type      Bind  Vis      Ndx Name
 15: 000000000201048      0 OBJECT LOCAL  DEFAULT 10 _GLOBAL_OFFSET_TABLE_

```

Symbol table shows that the GOT is at 201048. We can also see the same value in the section headers at '.got.plt'. That means we don't have to resolve symbols in order to get GOT's address which simplifies the first part. We can also learn that the GOT is 32 (0x20) bytes long in our example.

Note: the address of the GOT can also be found in the dynamic section for most executable files, under an entry tagged 'DT_PLTGOT.'

What's in the GOT?

We already know that

1. GOT+0x0 entry is unused
2. GOT+0x8 is an argument to second part
3. GOT+0x10 is function reference to second part

But what about the rest, starting at 201060 in our example? Here we have only one reference so it's obvious, but what if we have more references? How should we know which symbol is associated to which entry?

```
$ readelf -a test
Section Headers:
 [Nr] Name                Type          Address              Offset
     Size                EntSize          Flags Link Info Align
  [ 4] '.rela.plt'          RELA           0000000002001e8      000001e8
      000000000000018      000000000000018      AI      2    10    8

Relocation section '.rela.plt' at offset 0x1e8 contains 1 entries:
Offset             Info             Type              Sym. Value      Sym. Name + Addend
000000201060      000100000007      R_X86_64_JUMP_SLO 0000000000000000 printf + 0
```

How convenient that another table is also recorded in the section headers. It's called '.rela.plt' and describes exactly that. When the resolver receives control, the second parameter on the stack will be an index into this section, which is the offset (in 16-byte entries on x86-64, or 8-byte entries on 32-bit x86) of an entry like this that points to the symbol being requested.

Where are my libraries?

So far we assumed that shared libraries are already loaded. It's the case with libc, but how do we know what other shared libraries the executable wants?

```
$ readelf -a test
Section Headers:
 [Nr] Name                Type          Address              Offset
     Size                EntSize          Flags Link Info Align
  [ 6] '.dynamic'          DYNAMIC       000000000200220      00000220
      000000000000110      000000000000010      WA      3     0     8

Dynamic section at offset 0x220 contains 12 entries:
Tag              Type              Name/Value
0x0000000000000001 (NEEDED)             Shared library: [libc.so]
```

Not surprising that the answer lies in the section header again. There's a table pointer called `'.dynamic'`. That table has several records, but what we really are interested in is the ones marked by `"NEEDED"`.

Symbol look up

To find out `printf`'s address we should locate its symbol first in the shared library.

```
$ readelf -a libc.so
Section Headers:
 [Nr] Name                Type          Address              Offset
      Size              EntSize          Flags  Link  Info  Align
 [ 2]  '.dynsym'            DYNSYM         0000000100000218     00000218
      00000000000003c0    0000000000000018    A      3      1      8
 [ 3]  '.dynstr'            STRTAB         00000001000005d8     000005d8
      0000000000000124    0000000000000000    A      0      0      1

Symbol table '.dynsym' contains 40 entries:
 Num:  Value              Size Type          Bind  Vis      Ndx Name
  8:  0000000100000175      93 FUNC          GLOBAL DEFAULT  1  printf
```

Bingo! It is 100000175 in our example.

Gimme code!

I've put all the above together in a very simple example, see `elftool.c` (<https://gitlab.com/bztsrc/osz/-/blob/master/tools/elftool.c>) on gitlab.

When I run it on the executable it gives:

```
$ gcc elftool.c -o elftool
$ ./elftool -d mytestelf.o
Stringtable 000003f0 (118 bytes), symbols 000002b8 (312 bytes, one entry 24)

--- IMPORT ---
Dynamic 00001e20 (464 bytes, one entry 16):
 0. /lib/libc.so.6

GOT 00002000 (104 bytes), Rela 000004d0 (240 bytes, one entry 24):
 0. 00602018 +0 puts
 1. 00602020 +0 fread
 2. 00602028 +0 fclose
 3. 00602030 +0 printf
 4. 00602038 +0 strcmp
 5. 00602040 +0 ftell
 6. 00602048 +0 malloc
 7. 00602050 +0 fseek
 8. 00602058 +0 fopen
 9. 00602060 +0 exit

--- EXPORT ---
```

As you can see here we have an import section, but nothing to be exported. Now let's see a shared library!

```
$ ./elftool -d /lib/libc.so.6
Stringtable 00011038 (23041 bytes), symbols 00003d90 (53928 bytes, one entry 24)

--- IMPORT ---
Dynamic 00197b60 (496 bytes, one entry 16):
 0. /lib/ld-linux-x86-64.so.2

GOT 00198000 (88 bytes), Rela 0001f760 (192 bytes, one entry 24):
```

```
0. 00398050 +844d0
1. 00398048 +a8560
2. 00398040 +7ff80
3. 00398038 +867c0
4. 00398030 +823f0
5. 00398028 +82780
6. 00398020 +a8640
7. 00398018 +83b50

--- EXPORT ---
8. 0008e850 __strspn_c1
9. 0006aad0 putwchar
10. 000f8640 __gethostname_chk
11. 0008e870 __strspn_c2
12. 0010f210 setrpcent
13. 0009eda0 __wcstod_l
14. 0008e8a0 __strspn_c3
15. 000e8d10 epoll_create
16. 000d1b50 sched_get_priority_min
17. 000f8660 __getdomainname_chk
18. 000e8f20 klogctl
19. 0002c380 __tolower_l
20. 0004f440 dprintf
21. 000b8e00 setuid
22. 000a3d20 __wcscoll_l
... lot more lines to come ...
```

This time it has hell a lot of functions to export, and also it imports the dynamic linker of Linux with addend offsets in the GOT.

Summary

To summarize a dynamic linker should be look like:

1. load-time linker
 - 1.1. locates GOT by '.got.plt' section, and saves that address in GOT+0x8
 - 1.2. stores second part's address at GOT+0x10
 - 1.3. reads '.dynamic' section to load shared libraries
2. run-time reference resolver
 - 2.1. it's called by a helper
 - 2.2. that helper places index-0x18 and the address of GOT as arguments on the stack
 - 2.3. it has to locate '.rela.plt' to get the symbol for the reference
 - 2.4. the symbol is looked up in the shared library's '.dynsym' section to get relocated address
 - 2.5. that relocated address has to be saved in the GOT (index and base on the stack)
 - 2.6. clean up stack, restore registers and jump to the relocated address

That's all, hope it helps somebody! Good luck with implementing your own dynamic linker!

Retrieved from "https://osdev.wiki/wiki/Dynamic_Linkers"