



How kernel, compiler, and C library work together

Contents

[Kernel](#)

[C Library](#)

[Compiler / Assembler](#)

[Linker](#)

[Static Linking](#)

[Dynamic Linking](#)

[Shared Libraries](#)

[ABI - Application Binary Interface](#)

[Unresolved Symbols](#)

Kernel

The kernel is the core of an operating system. In a traditional design, it is responsible for memory management, I/O, interrupt handling, and various other things. And even while some modern designs like Microkernels or Exokernels move several of these services into user space, this matters little in the scope of this document.

The kernel makes its services available through a set of system calls; how they are called and what they do exactly differs from kernel to kernel.

C Library

Main articles: [C Library](#) and [Creating a C Library](#)

One thing up front: When you begin working on your kernel, you do not have a C library available. You have to provide everything yourself, except a few pieces provided by the compiler itself. You will also have to port an existing C library or write one yourself.

The C library implements the standard C functions (i.e., the things declared in `<stdlib.h>`, `<math.h>`, `<stdio.h>` etc.) and provides them in binary form suitable for linking with user-space applications.

In addition to standard C functions (as defined in the ISO standard), a C library might (and usually does) implement further functionality, which might or might not be defined by some standard. The standard C library says nothing about networking, for example. For Unix-like systems, the POSIX standard defines what is expected from a C library; other systems might differ fundamentally.

It should be noted that, in order to implement its functionality, the C library must call kernel functions. So, for your own OS, you can of course take a ready-made C library and just recompile it for your OS - but that requires that you tell the library how to call your kernel functions, and your kernel to actually provide those functions.

A more elaborate example is available in [Library Calls](#) or, you can use an existing [C Library](#) or [create your own C Library](#).

Compiler / Assembler

An Assembler takes (plaintext) source code and turns it into (binary) machine code; more precisely, it turns the source into *object* code, which contains additional information like symbol names, relocation information etc.

A compiler takes higher-level language source code, and either directly turns it into object code, or (as is the case with GCC) turns it into Assembler source code and invokes an Assembler for the final step.

The resulting object code does *not* yet contain any code for standard functions called. If you included e.g. `<stdio.h>` and used `printf()`, the object code will merely contain a *reference* stating that a function named `printf()` (and taking a `const char *` and a number of unnamed arguments as parameters) must be linked to the object code in order to receive a complete executable.

Some compilers use standard library functions *internally*, which might result in object files referencing e.g. `memset()` or `memcpy()` even though you did not include the header or used a function of this name. You will have to provide an implementation of these functions to the linker, or the linking will fail. The GCC freestanding environment expects only the functions `memset()`, `memcpy()`, `memcmp()`, and `memmove()`, as well as the [libgcc](#) library. Some advanced operations (e.g. 64-bits divisions on a 32-bits system) might involve *compiler-internal* functions. For [GCC](#), those functions are residing in `libgcc`. The content of this library is agnostic of what OS you use, and it won't taint your compiled kernel with licensing issues of whatever sort.

Linker

A linker takes the object code generated by the compiler / assembler, and *links* it against the C library (and / or `libgcc.a` or whatever link library you provide). This can be done in two ways: static, and dynamic.

Static Linking

When linking statically, the linker is invoked during the build process, just after the compiler / assembler run. It takes the object code, checks it for unresolved references, and checks if it can resolve these references from the available libraries. It then adds the binary code from these libraries to the executable; after this process, the executable is *complete*, i.e. when running it does not require anything but the kernel to be present.

On the downside, the executable can become quite large, and code from the libraries is duplicated over and over, both on disk and in memory.

Dynamic Linking

When linking dynamically, the linker is invoked during the *loading* of an executable. The unresolved references in the object code are resolved against the libraries currently present in the system. This makes the on-disk executable much smaller, and allows for in-memory space-saving strategies such as *shared libraries* (see below).

On the downside, the executable becomes dependent on the presence of the libraries it references; if a system does not have those libraries, the executable cannot run.

Shared Libraries

A popular strategy is to *share* dynamically linked libraries across multiple executables. This means that, instead of attaching the binary of the library to the executable image, the references in the executable are tweaked, so that all executables refer to the same in-memory representation of the required library.

This requires some trickery. For one, the library must either not have any *state* (static or global data) at all, or it must provide a separate *state* for each executable. This gets even trickier with multi-threaded systems, where one executable might have more than one simultaneous control flow.

Second, in a virtual memory environment, it is usually impossible to provide a library to all executables in the system at the same virtual memory address. To access library code at an arbitrary virtual address requires the library code to be *position independent* (which can be achieved e.g. by setting the `-PIC` command line option for the GCC compiler). This requires support of the feature by the binary format (relocation tables), and can result in slightly less efficient code on some architectures.

ABI - Application Binary Interface

The ABI of a system defines how library function calls and kernel system calls are actually done. This includes whether parameters are passed on the stack or in registers, how function entry points are located in libraries, and other such concerns.

When using static linkage, the resulting executable depends on the kernel using the same ABI as the one the executable was built for; when using dynamic linkage, the executable depends on the libraries' ABI staying the same.

Unresolved Symbols

The linker is the stage where you will find out about stuff that has been added without your knowledge, and which is not provided by your environment. This can include references to `alloca()`, `memcpy()`, or several others. This is usually a sign that either your toolchain or your command line options are not correctly set up for compiling your own OS kernel - or that you are

using functionality that is not yet implemented in your C library / runtime environment! You will most certainly run into trouble if you are not using a cross-compiler and the libgcc library and have implementations of memcpy, memmove, memset and memcmp.

Other symbols, such as __udiv* or __builtin_saveregs, are available in libgcc. If you get errors about missing such symbols, remember that you need to link with libgcc.

Retrieved from "https://osdev.wiki/wiki/How_kernel,_compiler,_and_C_library_work_together"

Content is available under CC0 Public Domain unless otherwise noted.